


# NLP Course

## Quantization – QLoRA NeurIPS LLM Efficiency Challenge

Nguyen Quoc Thai



# CONTENT

<b>1</b>	<b>Quantization</b>
<b>2</b>	<b>QLoRA</b>
<b>3</b>	<b>NeurIPS LLM Efficiency Challenge</b>

# 1 – Quantization



## Floating Point Number

- A floating point number is a positive or negative whole number with a decimal point

*man* → 

0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
-----	------	-----	-----	------	------	------

*woman* → 

0.7	0.3	0.9	-0.7	0.1	-0.5	-0.4
-----	-----	-----	------	-----	------	------

*king* → 

0.5	-0.4	0.7	0.8	0.9	-0.7	-0.6
-----	------	-----	-----	-----	------	------

*queen* → 

0.8	-0.1	0.8	-0.9	0.8	-0.5	-0.9
-----	------	-----	------	-----	------	------

# 1 – Quantization



## Tensor

➤ Tensor: multidimensional array

```
import torch  
data = torch.rand(3,3)  
data
```

```
tensor([[0.5123, 0.6089, 0.1713],  
        [0.2419, 0.8776, 0.8224],  
        [0.8413, 0.7830, 0.7792]])
```

# 1 – Quantization



## Tensor Properties

- Shape
- Device: CPU (-1), GPU (Cuda:0,...)
- Data Type

```
data.dtype
```

```
torch.float32
```

```
data.shape
```

```
torch.Size([3, 3])
```

```
data.get_device()
```

```
-1
```

Data type	dtype	CPU tensor	GPU tensor
32-bit floating point	<code>torch.float32</code> or <code>torch.float</code>	<code>torch.FloatTensor</code>	<code>torch.cuda.FloatTensor</code>
64-bit floating point	<code>torch.float64</code> or <code>torch.double</code>	<code>torch.DoubleTensor</code>	<code>torch.cuda.DoubleTensor</code>
16-bit floating point	<code>torch.float16</code> or <code>torch.half</code>	<code>torch.HalfTensor</code>	<code>torch.cuda.HalfTensor</code>
8-bit integer (unsigned)	<code>torch.uint8</code>	<code>torch.ByteTensor</code>	<code>torch.cuda.ByteTensor</code>
8-bit integer (signed)	<code>torch.int8</code>	<code>torch.CharTensor</code>	<code>torch.cuda.CharTensor</code>
16-bit integer (signed)	<code>torch.int16</code> or <code>torch.short</code>	<code>torch.ShortTensor</code>	<code>torch.cuda.ShortTensor</code>
32-bit integer (signed)	<code>torch.int32</code> or <code>torch.int</code>	<code>torch.IntTensor</code>	<code>torch.cuda.IntTensor</code>
64-bit integer (signed)	<code>torch.int64</code> or <code>torch.long</code>	<code>torch.LongTensor</code>	<code>torch.cuda.LongTensor</code>
Boolean	<code>torch.bool</code>	<code>torch.BoolTensor</code>	<code>torch.cuda.BoolTensor</code>

# 1 – Quantization



## FP32: Single Precision Floating Point

- 1 bit sign
  - 8 bits exponent
  - 23 bits fraction (precision)
- FP32: default => Weights, activations and other values in Neural Networks

$$+0.15625 = (+1) \times 2^{-3} \times 1.25$$


Sign 1 bit

Exponent 8 bits

Base: 2 or 10

Precision 23 bits

FP32



# 1 – Quantization



## FP32: Single Precision Floating Point

➤ Backward Propagation

➤ FP32: default => Weights, activations and other values in Neural Networks

Use a data type with fewer bits to represent floating point

$$+0.15625 = (+1) \times 2^{-3} \times 1.25$$


Sign 1 bit


Exponent 8 bits

Base: 2 or 10

Precision 23 bits

FP32







# 1 – Quantization

## FP16: Half Precision Floating Point

- 1 bit sign
- 5 bits exponent
- 10 bits fraction (precision)









# 1 – Quantization


## BFLOAT16: Brain Floating Point

- 1 bit sign
- 8 bits exponent
- 7 bits fraction (precision)



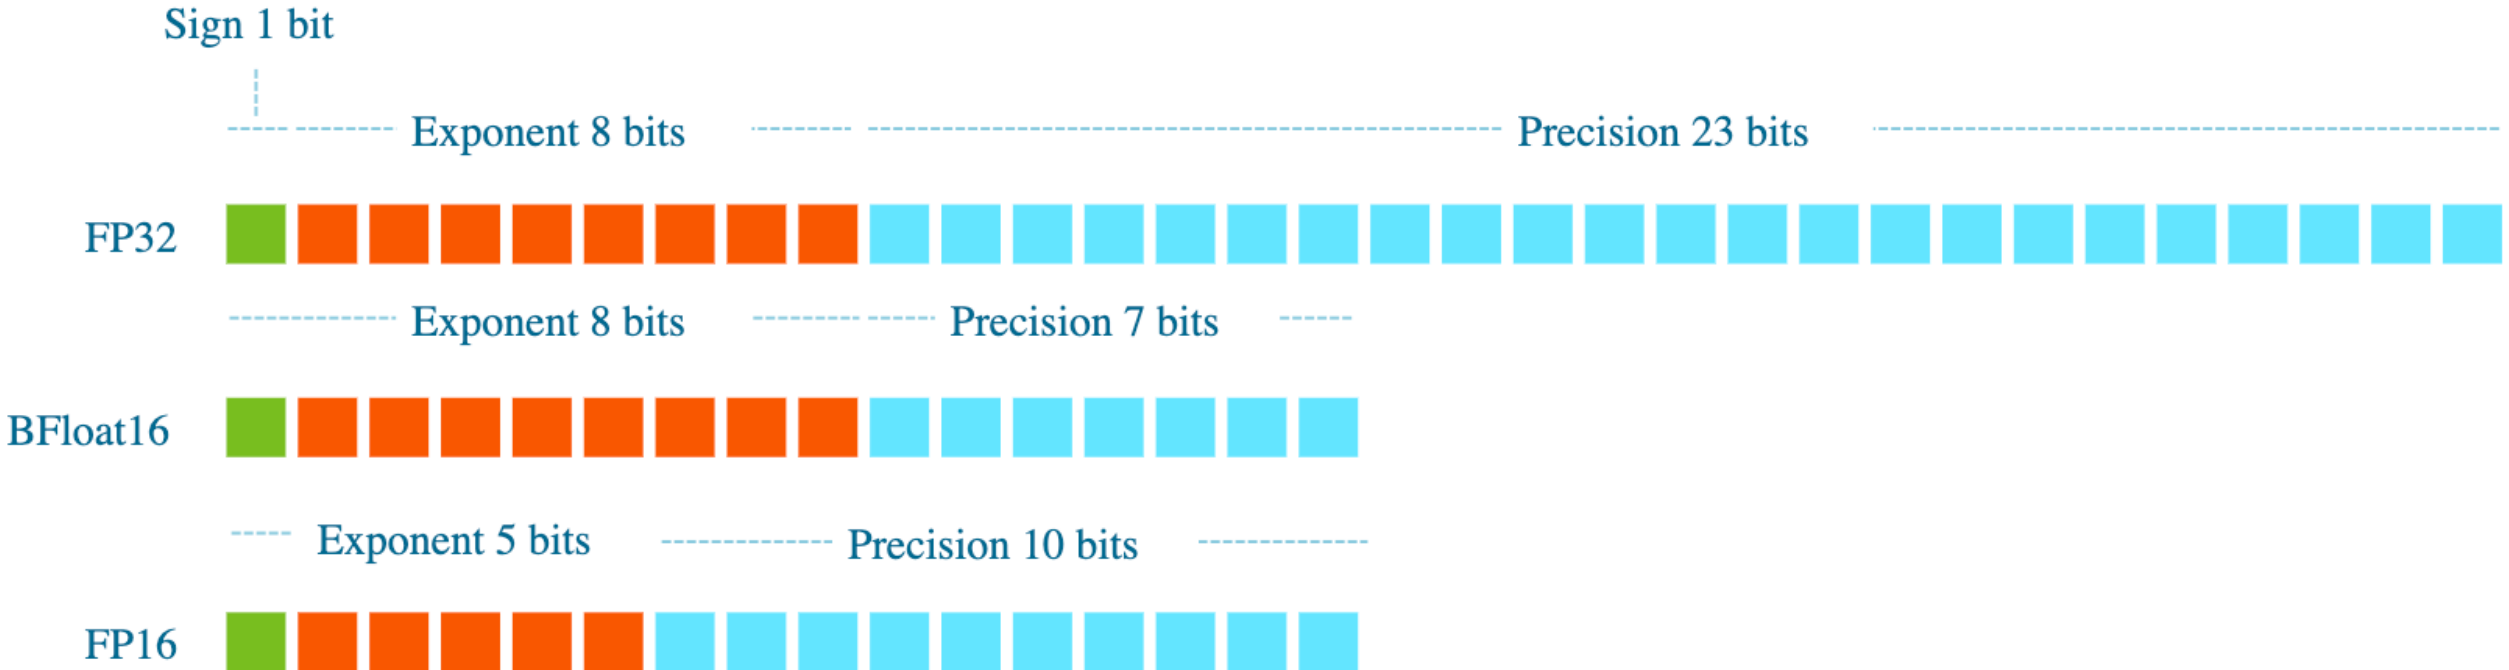


# 1 – Quantization



## Quantization

- Quantization: mapping input values from a large set (often a continuous set) to outputs values in a (countable) smaller set.
- Ex: Rounding and truncation



# 1 – Quantization




## Quantization

- Quantize from source dtype FP32 to target dtype INT8
- INT8: [-127, 127]

$$X^{\text{Int8}} = \text{round} \left( \frac{127}{\text{absmax}(X^{\text{FP32}})} X^{\text{FP32}} \right) = \text{round}(c^{\text{FP32}} X^{\text{FP32}})$$

c: constant



# 1 – Quantization



## Quantization

- Dequantize from target dtype INT8 to source dtype FP32
- INT8: [-127, 127]

$$\text{dequant}(c^{\text{FP32}} X^{\text{FP32}}) = \frac{X^{\text{Int8}}}{c^{\text{FP32}}} = X^{\text{FP32}}$$

$$C = \frac{127}{0.4} = 317.5$$




# 1 – Quantization





## Smaller and Faster

### Memory Usage



### Performance

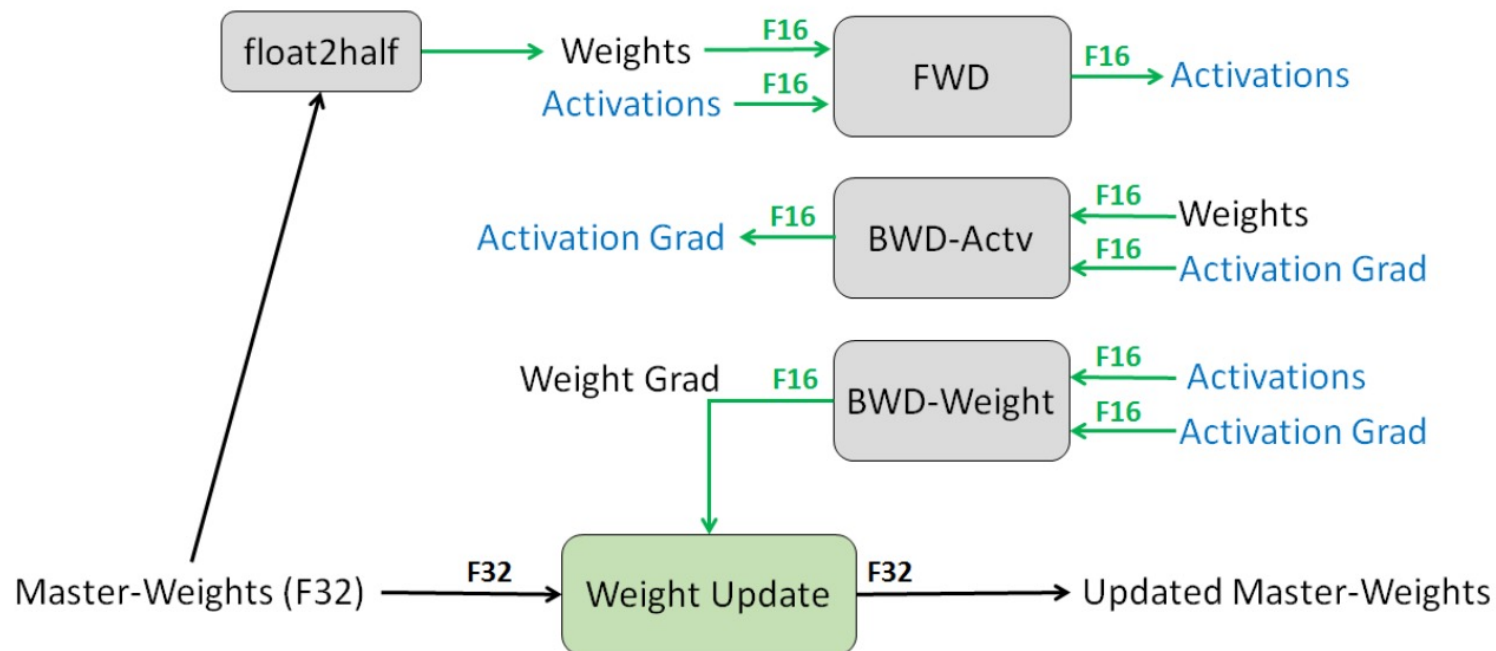


# 1 – Quantization



## Mixed Precision Training

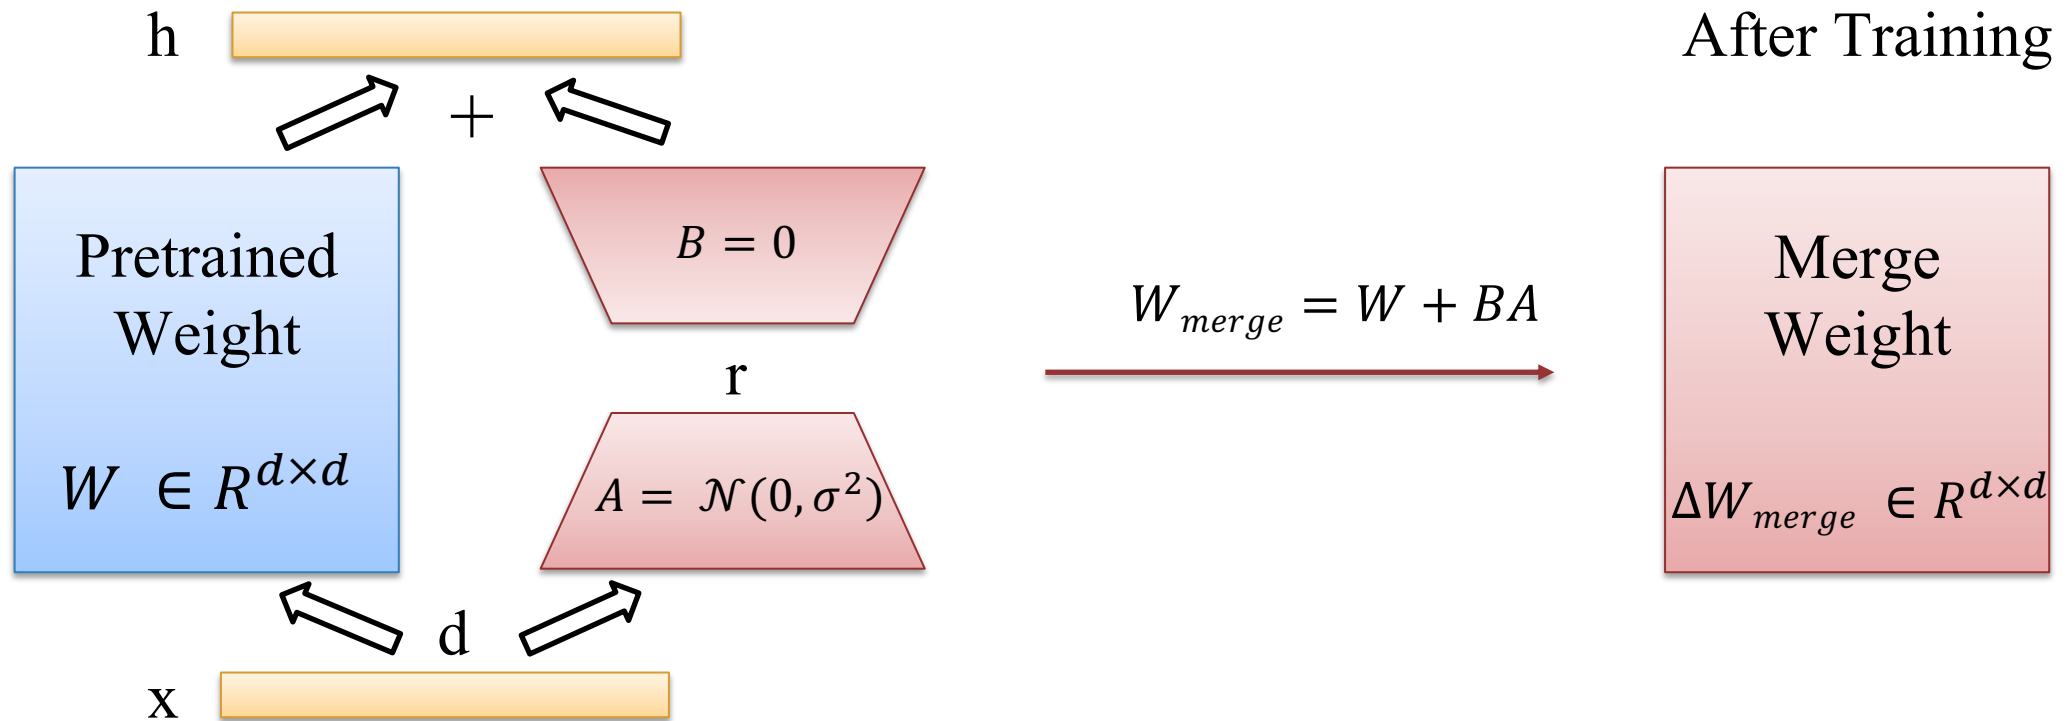
- Mixed Precision Training: Not a floating point data type but a method
- Use a combination of FP16 and FP32 to reduce the memory and math bandwidth






## LoRA: Low-Rank Adaptation

- Freezes the pretrained model weights and injects trainable rank decomposition matrices into each layer of the Transformer architecture





## LoRA: Low-Rank Adaptation

- LoRA can even outperform full finetuning training only 2% of the parameters


Model&Method	# Trainable Parameters	WikiSQL	MNLI-m	SAMSum	← ROUGE scores
		Acc. (%)	Acc. (%)	R1/R2/RL	
Full finetuning → GPT-3 (FT)	175,255.8M	<b>73.8</b>	89.5	52.0/28.0/44.5	
Only tune bias vectors → GPT-3 (BitFit)	14.2M	71.3	91.0	51.3/27.4/43.5	
Prompt tuning → GPT-3 (PreEmbed)	3.2M	63.1	88.6	48.3/24.2/40.5	
Prompt tuning → GPT-3 (PreLayer)	20.2M	70.1	89.5	50.8/27.3/43.5	
Prefix tuning → GPT-3 (Adapter <sup>H</sup> )	7.1M	71.9	89.8	53.0/28.9/44.8	
Prefix tuning → GPT-3 (Adapter <sup>H</sup> )	40.1M	73.2	<b>91.5</b>	53.2/29.0/45.1	
GPT-3 (LoRA)	4.7M	73.4	<b>91.7</b>	<b>53.8/29.8/45.9</b>	
GPT-3 (LoRA)	37.7M	<b>74.0</b>	<b>91.6</b>	53.4/29.2/45.1	





## Challenges of Quantization Method

- Information Loss
- Example: quantize from INT3 to INT2



Speed



Performance






## Challenges of Quantization Method

➤ Linear Quantization (Ignore the distribution on the source data type)


➤ Example: 
$$X^{\text{Int8}} = \text{round} \left( \frac{127}{\text{absmax}(X^{\text{FP32}})} X^{\text{FP32}} \right) = \text{round}(c^{\text{FP32}} X^{\text{FP32}})$$





## Challenges of Quantization Method

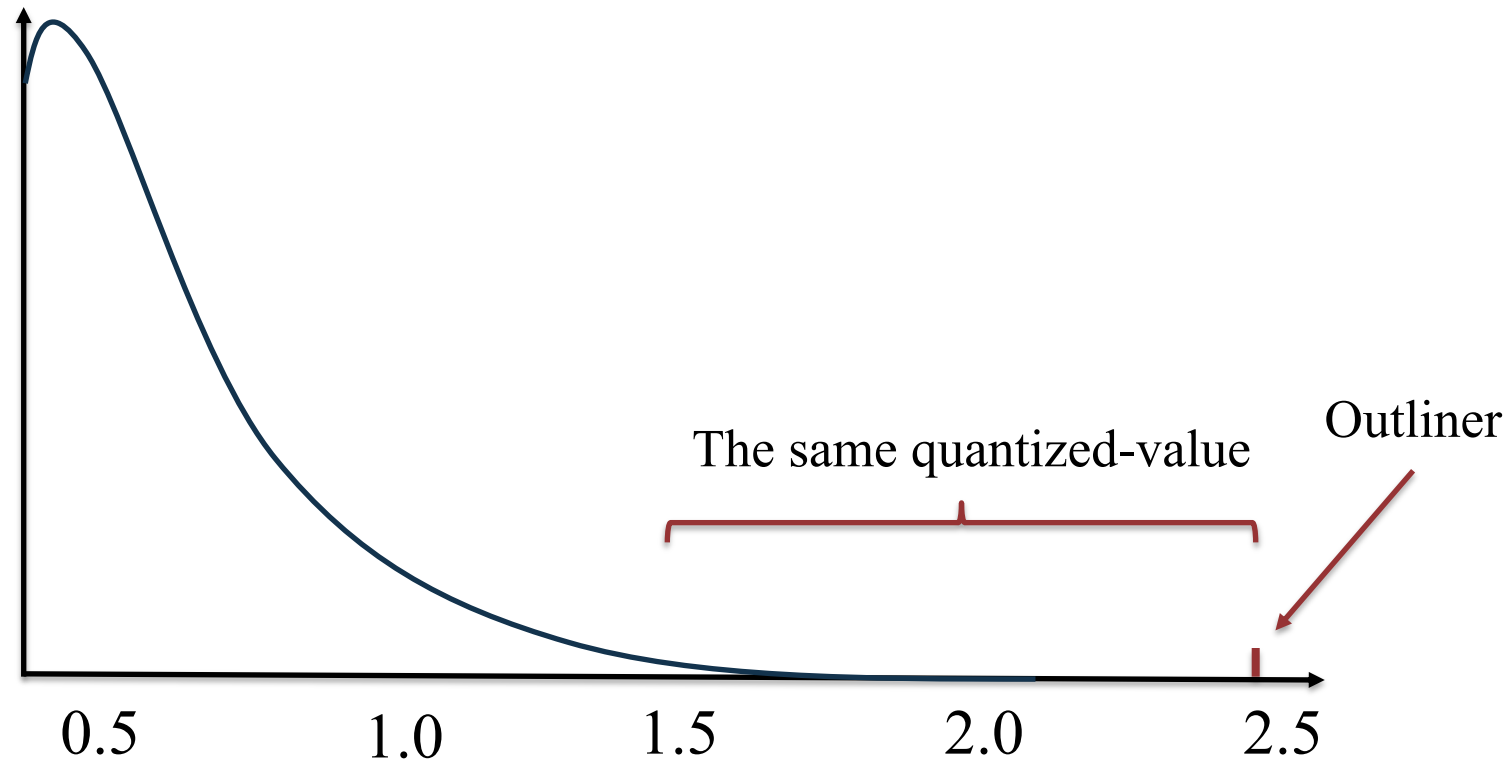
- Quantile Quantization
- Quantiles: cut points dividing the range of a probability distribution into continuous intervals with equal probabilities





## Challenges of Quantization Method


- Outliner in Quantization: appear few times but is far away from other values
- Outliner often very important (Attention score)





## Challenges of Quantization Method

- Block-wise Quantization: split a tensor into many chunks, quantize individual chunks
- Block size: number of elements in a chunk





## QLoRA: Efficient Finetuning of Quantized LLMs

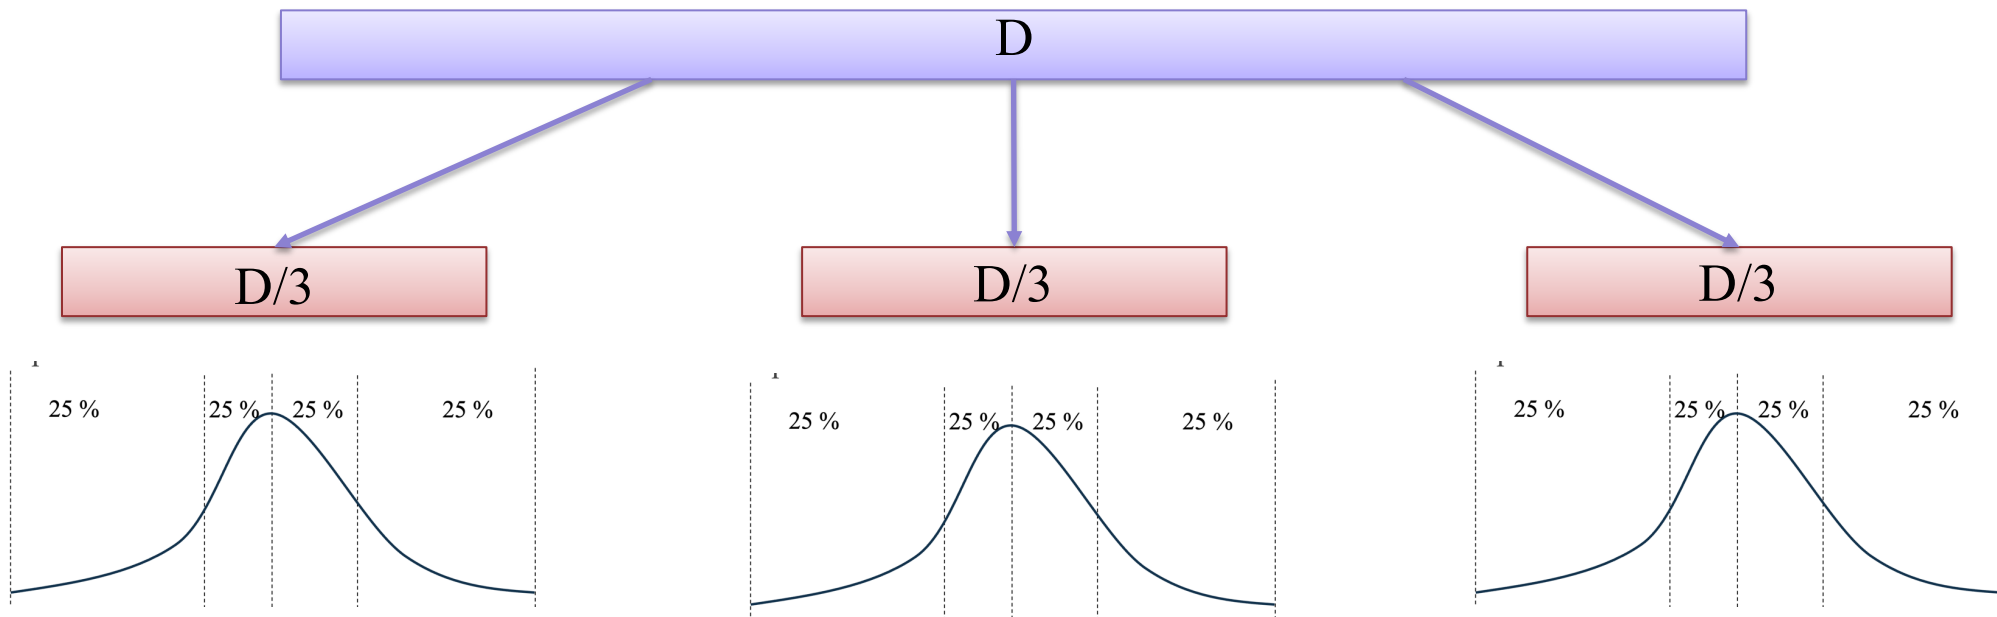
**QLoRA: save memory without sacrificing performance**

- 4-bit NormalFloat (NF4) via Block-wise Quantization
- Double Quantization
- Paged Optimizers
- Combined with LoRA



## 4-bit NormalFloat (NF4)

- Step 1: Find quantiles in each chunks




=> The main limitation of quantile quantization: process of quantile estimation very expensive



### 4-bit NormalFloat (NF4)

- Step 1: Find quantiles in each chunks
- Use fixed distribution: zero-mean normal distribution with standard deviation  $\sigma$







### 4-bit NormalFloat (NF4)

- Step 1: estimate the  $2^k + 1$  quantiles of a theoretical  $N(0, 1)$  distribution to obtain a  $k$ -bit quantile quantization data type for normal distributions as follows:

$$q_i = \frac{1}{2} \left( Q_X \left( \frac{i}{2^k + 1} \right) + Q_X \left( \frac{i + 1}{2^k + 1} \right) \right)$$

QX: the quantile function of the standard normal distribution  $N(0,1)$




### 4-bit NormalFloat (NF4)

- Step 1: Estimate the  $2^k + 1$  quantiles of a theoretical  $N(0, 1)$  distribution to obtain a k-bit quantile quantization data type for normal distributions
- Step 2: Take this data type and normalize its values into the  $[-1, 1]$  range
- Step 3: Quantize an input weight tensor by normalizing it into the  $[-1, 1]$  range through absolute maximum rescaling



### 4-bit NormalFloat (NF4)

- Problem: Quantiles not have an exact representation of zero (Symmetric)
- Important property to quantize padding and other zero-valued elements with no error






### 4-bit NormalFloat (NF4)

Solution: create an asymmetric data type by estimating quantiles  $q_i$  of two range:


- $2^{k-1}$  for the negative part
- $2^{k-1} + 1$  for the positive part
- Then unify these sets of  $q_i$  and remove one of the two zeros that occurs in both sets

=> **K-bit NormalFloat (NFk) data type**




## 4-bit NormalFloat (NF4)

- Use 4 bits to representation
- Normalize into  $[-1, 1]$  range
- An asymmetric data type: an exact representation of zero
- Quantiles based on zero-mean normal distribution with standard deviation  $\sigma$



## ! Double Quantization

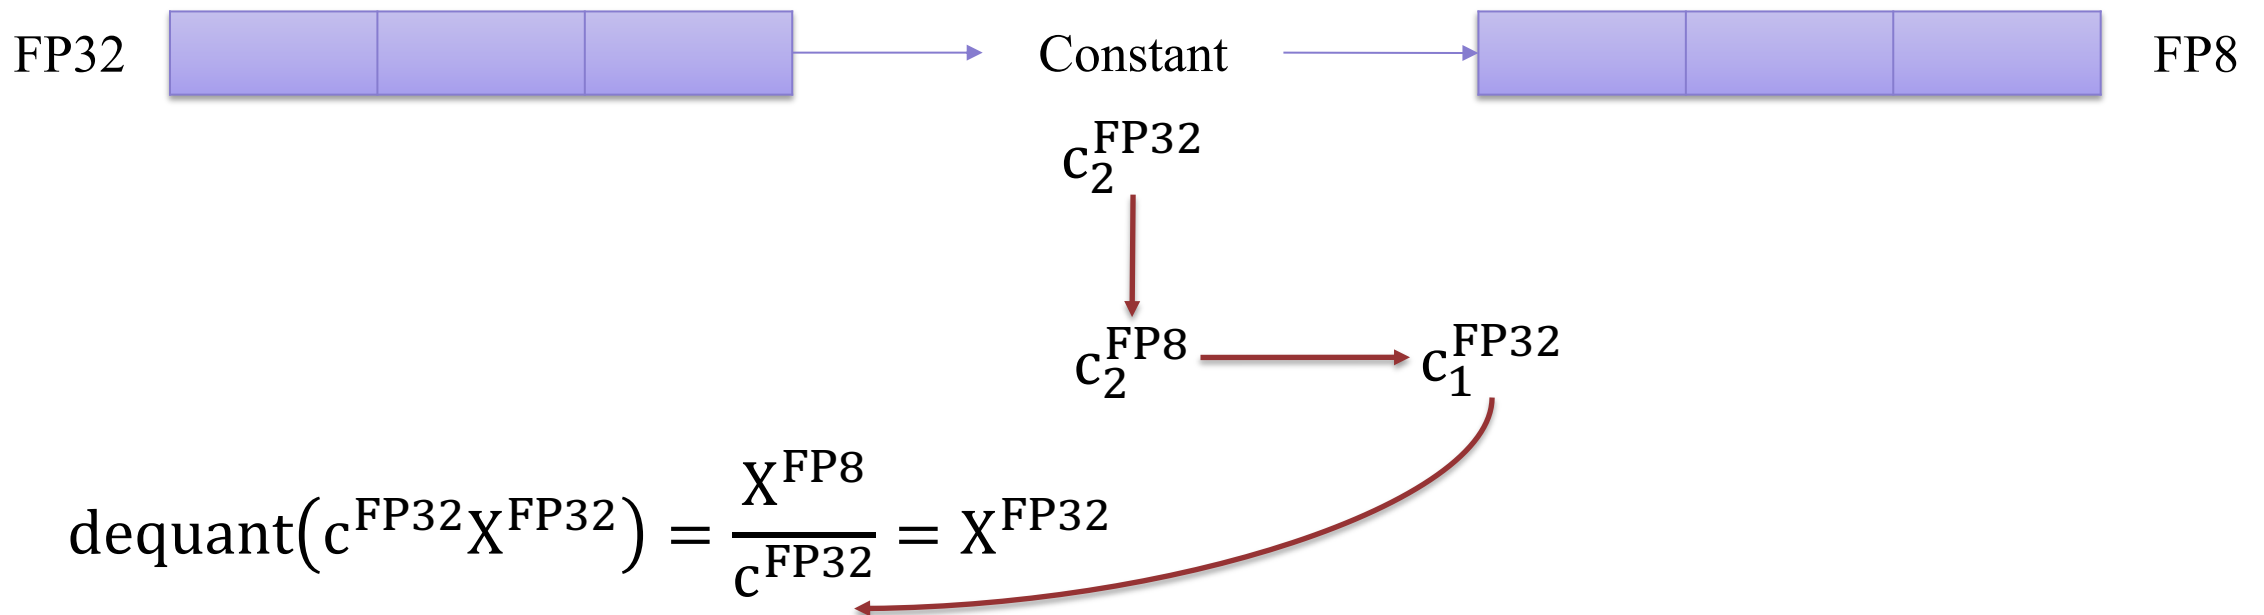




## Double Quantization

- The process of quantizing the quantization constants for additional memory savings


Chunk





## Double Quantization

- The process of quantizing the quantization constants for additional memory savings








## Paged Optimizers

- Page Optimizers to manage memory spikes
- Allocate paged memory for the optimizer states which are then automatically evicted to CPU RAM when the GPU runs out-of-memory and paged back into GPU memory when the memory is needed in the optimizer update step

## 2 – QLoRA



## QLoRA





## QLoRA

- Given a projection  $XW = Y, X \in \mathbb{R}^{b \times h}, W \in \mathbb{R}^{h \times o}$ , LoRA computes:

$$Y = XW + sXL_1L_2$$

$L_1 \in \mathbb{R}^{h \times r}, L_2 \in \mathbb{R}^{r \times o}$ ,  $s$  is a scalar



## QLoRA

➤ Quantized base model with a single LoRA adapter:

$$Y^{BF16} = X^{BF16} \text{doubleDequant}(c_1^{FP32}, c_2^{k\text{-bit}}, W^{NF4}) + X^{BF16} L_1^{BF16} L_2^{BF16}$$

$$\text{doubleDequant}(c_1^{FP32}, c_2^{k\text{-bit}}, W^{NF4}) = \text{dequant}(\text{dequant}(c_1^{FP32}, c_2^{k\text{-bit}}), W^{4\text{bit}}) = W^{BF16}$$

NF4 for W and FP8 for c2

A block size of 64 for W for higher quantization precision

A block size of 256 for c2 to conserve memory



## Result

Model / Dataset	Params	Model bits	Memory	ChatGPT vs Sys	Sys vs ChatGPT	Mean	95% CI
GPT-4	-	-	-	119.4%	110.1%	<b>114.5%</b>	2.6%
Bard	-	-	-	93.2%	96.4%	<b>94.8%</b>	4.1%
<b>Guanaco</b>	65B	4-bit	41 GB	96.7%	101.9%	<b>99.3%</b>	4.4%
Alpaca	65B	4-bit	41 GB	63.0%	77.9%	70.7%	4.3%
FLAN v2	65B	4-bit	41 GB	37.0%	59.6%	48.4%	4.6%
<b>Guanaco</b>	33B	4-bit	21 GB	96.5%	99.2%	<b>97.8%</b>	4.4%
Open Assistant	33B	16-bit	66 GB	91.2%	98.7%	94.9%	4.5%
Alpaca	33B	4-bit	21 GB	67.2%	79.7%	73.6%	4.2%
FLAN v2	33B	4-bit	21 GB	26.3%	49.7%	38.0%	3.9%
Vicuna	13B	16-bit	26 GB	91.2%	98.7%	<b>94.9%</b>	4.5%
<b>Guanaco</b>	13B	4-bit	10 GB	87.3%	93.4%	90.4%	5.2%
Alpaca	13B	4-bit	10 GB	63.8%	76.7%	69.4%	4.2%
HH-RLHF	13B	4-bit	10 GB	55.5%	69.1%	62.5%	4.7%
Unnatural Instr.	13B	4-bit	10 GB	50.6%	69.8%	60.5%	4.2%
Chip2	13B	4-bit	10 GB	49.2%	69.3%	59.5%	4.7%
Longform	13B	4-bit	10 GB	44.9%	62.0%	53.6%	5.2%
Self-Instruct	13B	4-bit	10 GB	38.0%	60.5%	49.1%	4.6%
FLAN v2	13B	4-bit	10 GB	32.4%	61.2%	47.0%	3.6%
<b>Guanaco</b>	7B	4-bit	5 GB	84.1%	89.8%	<b>87.0%</b>	5.4%
Alpaca	7B	4-bit	5 GB	57.3%	71.2%	64.4%	5.0%
FLAN v2	7B	4-bit	5 GB	33.3%	56.1%	44.8%	4.0%


# 3 – LLM Efficiency Challenge




## Challenge

**NeurIPS Large Language Model Efficiency Challenge:**  
**1 LLM + 1GPU + 1Day**

NeurIPS 2023 Challenge



# 3 – LLM Efficiency Challenge



## Challenge

### ➤ Approved Base Models

- Falcon
- LLaMA or Llama 2
- OpenLLaMA
- Red Pajama Base (not instruction tuned models)
- MPT
- OPT
- Bloom
- GPT Neo, J, NeoX, Pythia
- GPT2
- T5 (not Flan-T5)
- BART
- DeBERTa
- RoBERTa
- BERT
- ALBERT
- DistilBERT
- Electra
- UL2
- Cerebras (btlm, GPT)

### ➤ Approved Base Models

- Databricks-Dolly-15
- OpenAssistant Conversations Dataset (oasst1)
- The Flan Collection
- AllenAI Dolma
- RedPajama-Data-1T
- LIMA

# 3 – LLM Efficiency Challenge



## Methods and tools for efficient training on a single GPU

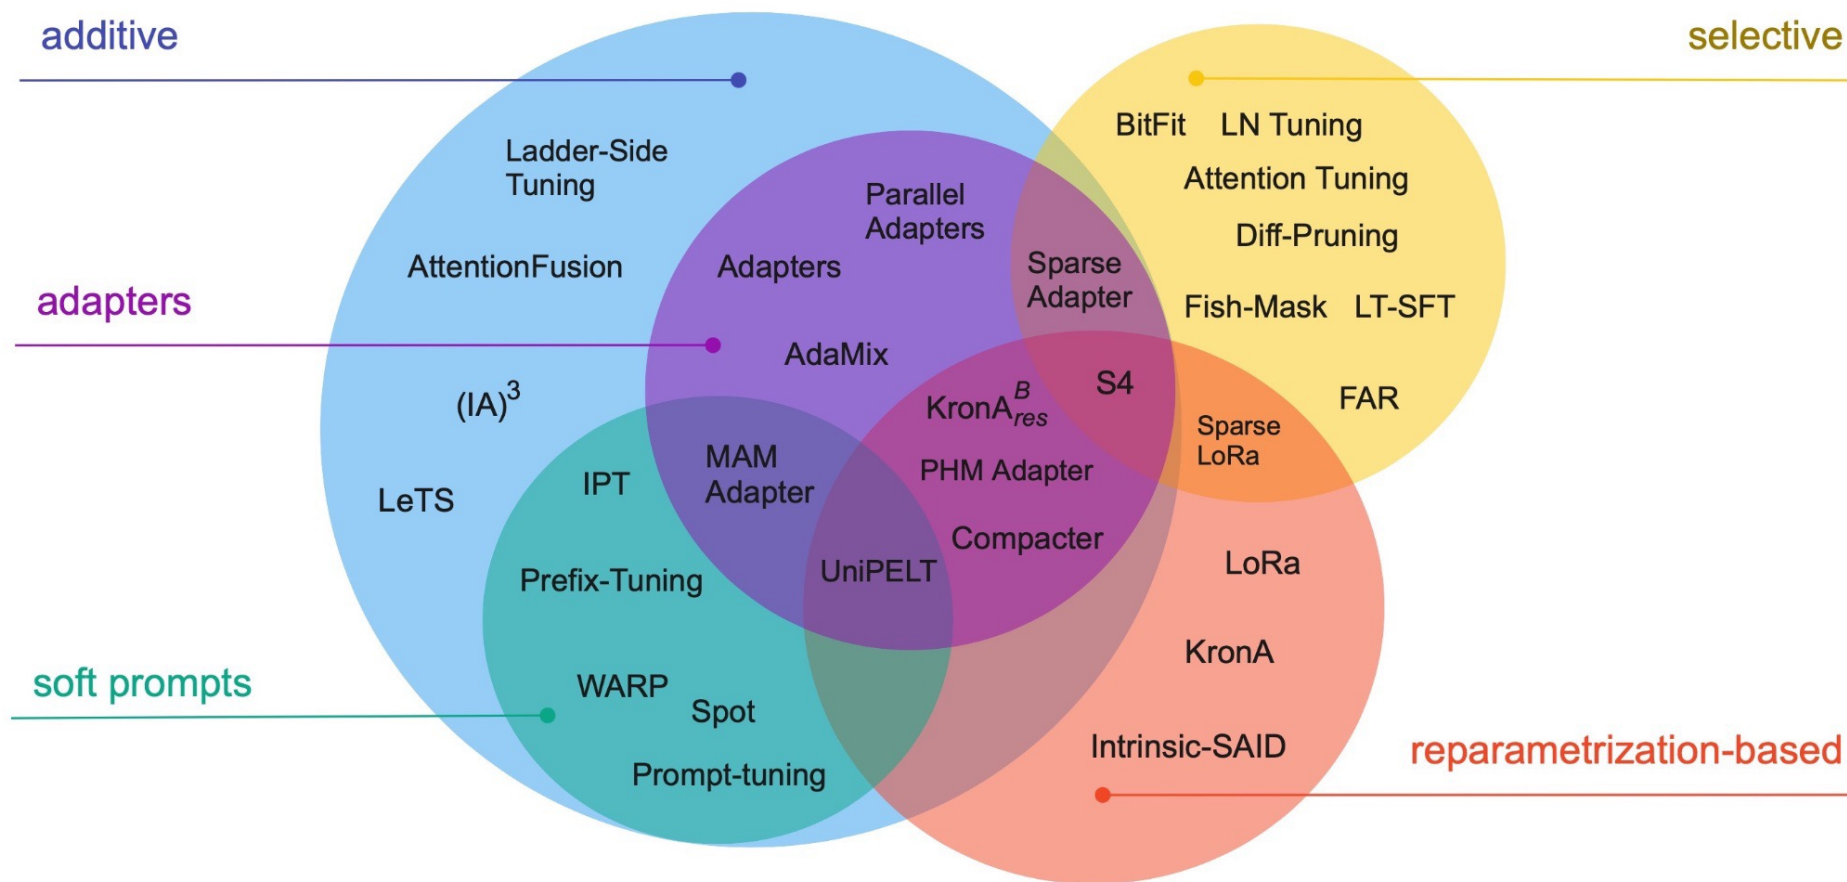
Method/tool	Improves training speed	Optimizes memory utilization
<u>Batch size choice</u>	Yes	Yes
<u>Gradient accumulation</u>	No	Yes
<u>Gradient checkpointing</u>	No	Yes
<u>Mixed precision training</u>	Yes	(No)
<u>Optimizer choice</u>	Yes	Yes
<u>Data preloading</u>	Yes	No
<u>DeepSpeed Zero</u>	No	Yes
<u>torch.compile</u>	Yes	No



# 3 – LLM Efficiency Challenge




## Parameter-Efficient Fine-Tuning



# 3 – LLM Efficiency Challenge



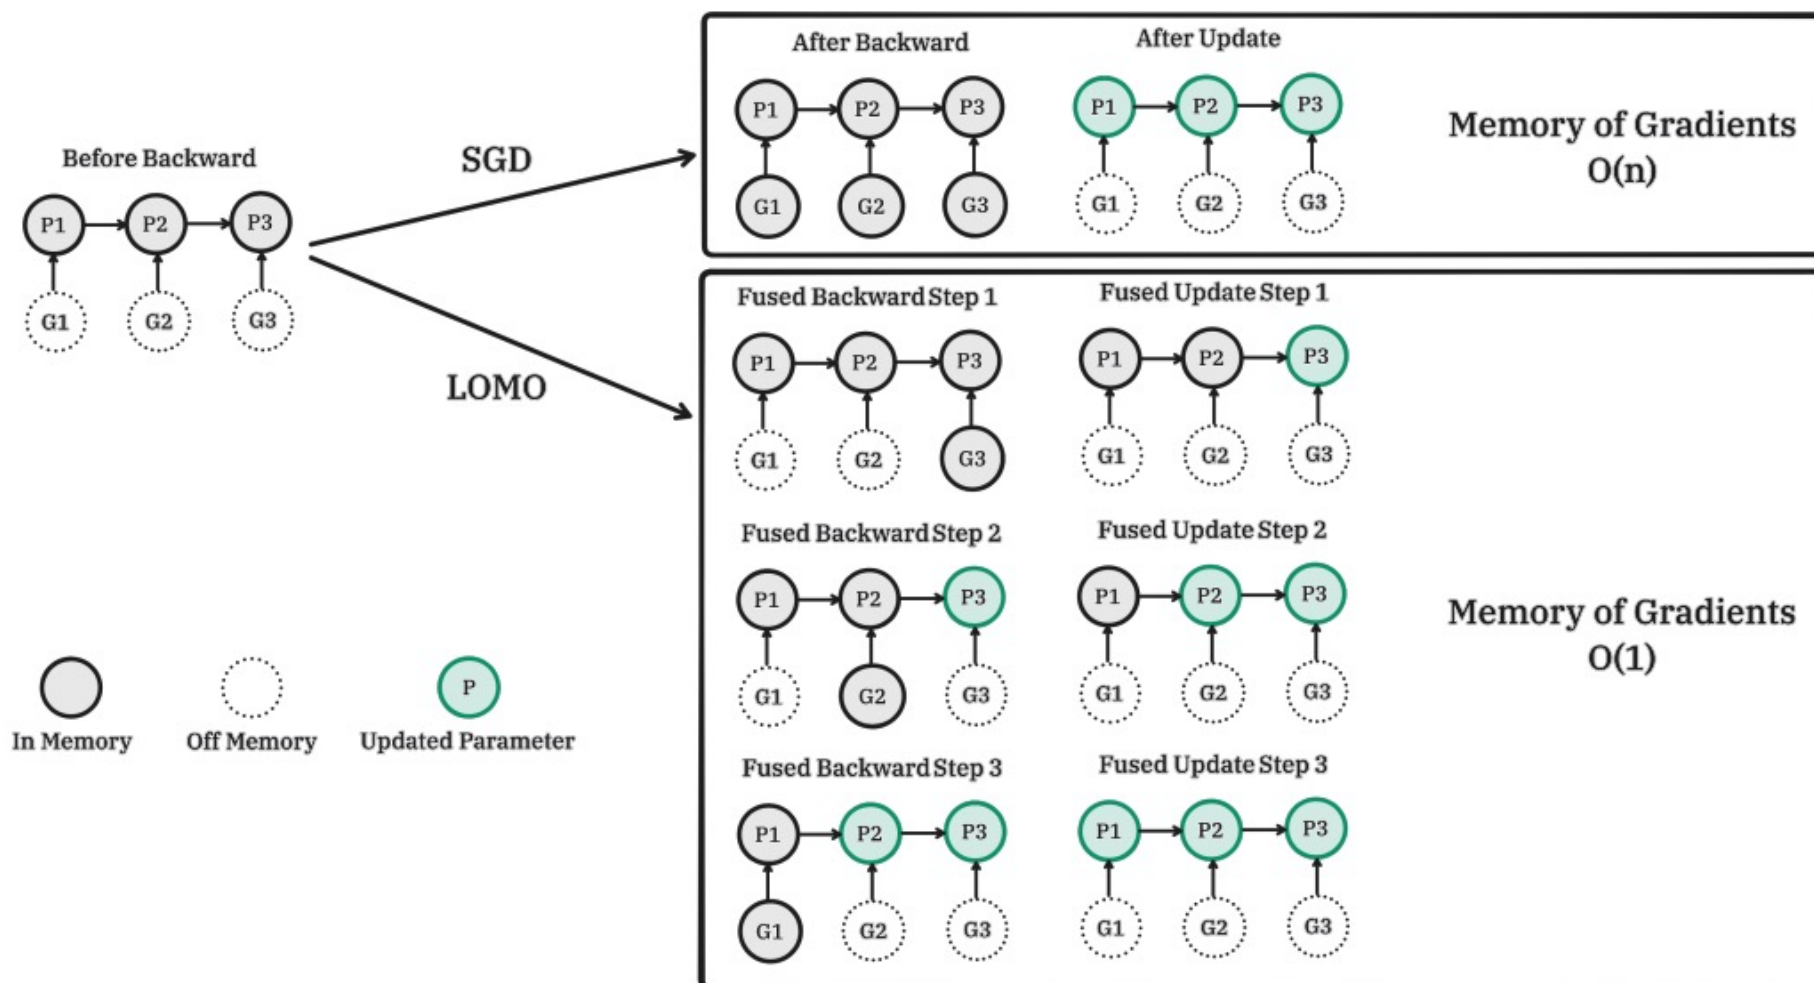
## Quatization




# 3 – LLM Efficiency Challenge



## Low-Memory Optimization (LOMO)





AI VIET NAM

@aivietnam.edu.vn

# 4 - Experiment



**Source code**



AI VIET NAM

@aivietnam.edu.vn

# Thanks!

## Any questions?